

- SSH command to connect to Hadoop EMR-
ssh -i /c/Users/Ankan\ Mazumdar/Downloads/emr-key-pair.pem
hadoop@ec2-18-220-141-223.us-east-2.compute.amazonaws.com

```
Ankan Mazumdar@DESKTOP-CMULEBA MINGW64 ~/Downloads
$ ssh -i emr_key_pair.pem hadoop@ec2-18-220-141-223.us-east-2.compute.amazonaws.com
Last login: Fri Sep 15 01:21:45 2023

  _| _|_ )
  _| (  /   Amazon Linux 2 AMI
 _|\_||_|

https://aws.amazon.com/amazon-linux-2/
23 package(s) needed for security, out of 57 available
Run "sudo yum update" to apply all updates.

EEEEEEEEEEEEEEEEEEEE MMMMMMMM          MMMMMMMM RRRRRRRRRRRRRRRR
E::::::::::::::::::::E M::::::::M          M::::::::M R:::::::::R
EE::::::::EEEEEEEEEEE M:::::::::M          M:::::::::M R::::RRRRRR::::R
  E::::E          EEEEE M:::::::::M          M:::::::::M RR::::R      R::::R
  E::::E          M:::::M M::::M M::::M M::::M          R:::R      R::::R
  E::::EEEEEEEEEEE M:::::M M::::M M::::M M::::M          R::RRRRRR::::R
  E::::::::::::::::E M:::::M M::::M M::::M M::::M          R:::::::::RR
  E::::EEEEEEEEEEE M:::::M M::::M M::::M M::::M          R::RRRRRR::::R
  E::::E          M:::::M M::::M M::::M M::::M          R:::R      R::::R
  E::::E          EEEEE M:::::M          MMM M:::::M          R:::R      R::::R
EE::::::::EEEEEEEEEEE M:::::M          M:::::M          R:::R      R::::R
E::::::::::::::::::::E M:::::M          M:::::M RR::::R      R::::R
EEEEEEEEEEEEEEEEEEEE MMMMMMMM          MMMMMMMM RRRRRRR      RRRRRR
```

- SCP file from local PC to the Hadoop master node-

```
Ankan Mazumdar@DESKTOP-CMULEBA MINGW64 ~/Downloads
scp -i emr-key-pair.pem ankanmazumdar.txt
hadoop@ec2-18-220-141-223.us-east-2.compute.amazonaws.com:/home/hadoop
```

```
Ankan Mazumdar@DESKTOP-CMULEBA MINGW64 ~/Downloads
$ scp -i emr_key_pair.pem ankanmazumdar.txt hadoop@ec2-18-220-141-223.us-east-2.compute.amazonaws.com:/home/hadoop/
ankanmazumdar.txt 100% 21 0.7KB/s 00:00
```

- Upload "myid.txt" in S3 bucket-

Amazon S3 > Buckets > a20541357-csp554 > A20541357.txt

A20541357.txt [Info](#)

[Copy S3 URI](#) [Download](#) [Open](#) [Object actions](#)

[Properties](#) | [Permissions](#) | [Versions](#)

Object overview

Owner 885e6b93b6399c4c404b6b4faec68d73484c78dc731875a7be31ba357d85bd0e	S3 URI s3://a20541357-csp554/A20541357.txt
AWS Region US East (Ohio) us-east-2	Amazon Resource Name (ARN) arn:aws:s3::a20541357-csp554/A20541357.txt
Last modified September 14, 2023, 20:42:52 (UTC-05:00)	Entity tag (Etag) fb6e6769f18d89f27576fdf8cccc9d4
Size 21.0 B	Object URL https://a20541357-csp554.s3.us-east-2.amazonaws.com/A20541357.txt
Type txt	
Key A20541357.txt	

9. (2 points) Execute the following `hdfs` command to list the files or directories that are listed (also indicating which is a file and which a directory):
- `hadoop dfs -ls /`

```
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -ls /
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

Found 5 items
drwxr-xr-x - hdfs hdfsadmingroup 0 2023-09-14 23:05 /apps
drwxr-xr-x - hadoop hdfsadmingroup 0 2023-09-15 01:45 /test
drwxrwxrwt - hdfs hdfsadmingroup 0 2023-09-14 23:09 /tmp
drwxr-xr-x - hdfs hdfsadmingroup 0 2023-09-14 23:05 /user
drwxr-xr-x - hdfs hdfsadmingroup 0 2023-09-14 23:05 /var
```

10. (2 points) Execute a command (you needed to figure out which one) to list the files and directories under the hdfs directory listed below: /user

`hadoop dfs -ls /user`

```
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -ls /user
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

Found 6 items
drwxrwxrwx - hadoop hdfsadmingroup 0 2023-09-14 23:05 /user/hadoop
drwxr-xr-x - mapred mapred 0 2023-09-14 23:05 /user/history
drwxrwxrwx - hdfs hdfsadmingroup 0 2023-09-14 23:05 /user/hive
drwxrwxrwx - hue hue 0 2023-09-14 23:05 /user/hue
drwxrwxrwx - oozie oozie 0 2023-09-14 23:09 /user/oozie
drwxrwxrwx - root hdfsadmingroup 0 2023-09-14 23:05 /user/root
```

11. (2 points) Execute a command to create the following HDFS directory: /user/csp554 Record the command you executed and include it in your assignment submission.

`hadoop dfs -mkdir -p /user/csp554`

```
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -mkdir -p /user/csp554/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -ls /user/csp554/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -ls /user/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

Found 7 items
drwxr-xr-x - hadoop hdfsadmingroup 0 2023-09-15 02:23 /user/csp554
drwxrwxrwx - hadoop hdfsadmingroup 0 2023-09-14 23:05 /user/hadoop
drwxr-xr-x - mapred mapred 0 2023-09-14 23:05 /user/history
drwxrwxrwx - hdfs hdfsadmingroup 0 2023-09-14 23:05 /user/hive
drwxrwxrwx - hue hue 0 2023-09-14 23:05 /user/hue
drwxrwxrwx - oozie oozie 0 2023-09-14 23:09 /user/oozie
drwxrwxrwx - root hdfsadmingroup 0 2023-09-14 23:05 /user/root
```

12. (2 points) Execute a command to create the following HDFS directory: /user/csp554-2 Record the command you executed and include it in your assignment submission.

`hadoop dfs -mkdir -p /user/csp554-2/`

```
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -mkdir -p /user/csp554-2/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -ls /user/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

Found 8 items
drwxr-xr-x - hadoop hdfsadmin 0 2023-09-15 02:23 /user/csp554
drwxr-xr-x - hadoop hdfsadmin 0 2023-09-15 02:24 /user/csp554-2
drwxrwxrwx - hadoop hdfsadmin 0 2023-09-14 23:05 /user/hadoop
drwxr-xr-x - mapred mapred 0 2023-09-14 23:05 /user/history
drwxrwxrwx - hdfs hdfsadmin 0 2023-09-14 23:05 /user/hive
drwxrwxrwx - hue hue 0 2023-09-14 23:05 /user/hue
drwxrwxrwx - oozie oozie 0 2023-09-14 23:09 /user/oozie
drwxrwxrwx - root hdfsadmin 0 2023-09-14 23:05 /user/root
```

13. (2 points) Execute a command that copies a given local file to the given hdfs directory : Source local file: /home/hadoop/myname.txt (where the actual name is your name as described above) Destination HDFS directory: /user/csp554 Record the command you executed and include it in your assignment submission.

`hadoop dfs -copyFromLocal -f /home/hadoop/ankanmazumdar.txt /user/csp554/`

```
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -copyFromLocal -f /home/hadoop/ankanmazumdar.txt /user/csp554/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -ls /user/csp554/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

Found 1 items
-rw-r--r-- 1 hadoop hdfsadmin 21 2023-09-15 02:35 /user/csp554/ankanmazumdar.txt
```

14. (2 points) Copy a file from one hdfs directory to another hdfs directory and write down the command Source hdfs file: /user/csp554/myname.txt (where the actual name is your name as described above) Destination HDFS directory: /user/csp554-2 Record the command you executed and include it in your assignment submission.

`hadoop dfs -cp /user/csp554/ankanmazumdar.txt /user/csp554-2/`

```
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -cp /user/csp554/ankanmazumdar.txt /user/csp554-2/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -ls /user/csp554-2/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

Found 1 items
-rw-r--r-- 1 hadoop hdfsadmin 21 2023-09-15 02:36 /user/csp554-2/ankanmazumdar.txt
```

15. (2 points) Copy the object myid.txt you uploaded to an S3 bucket into the Hadoop master node Linux file system.

```
aws s3 cp s3://a20541357-csp554/A20541357.txt /home/hadoop/A20541357.txt
```

```
[hadoop@ip-172-31-35-150 ~]$ aws s3 cp s3://a20541357-csp554/A20541357.txt /home/hadoop/A20541357.txt
download: s3://a20541357-csp554/A20541357.txt to ./A20541357.txt
[hadoop@ip-172-31-35-150 ~]$ ls /home/hadoop
A20541357.txt  ankanmazumdar.txt
[hadoop@ip-172-31-35-150 ~]$ ls -l /home/hadoop
total 8
-rw-rw-r-- 1 hadoop hadoop 21 Sep 15 01:42 A20541357.txt
-rw-r--r-- 1 hadoop hadoop 21 Sep 15 01:34 ankanmazumdar.txt
```

16. (2 points) Copy the same object myid.txt you created in an S3 bucket into HDFS into the directory /user/csp554-2. Note, the three slashes after the "hdfs:." After you executed the above command, execute another command (you needed to figure out which one) to list the files and directories under the hdfs directory listed below: /user/csp554-2. Write down the command you executed and also take a screen snapshot of names of the files or directories that are listed and include it in your assignment submission.

```
hadoop fs -cp s3://a20541357-csp554/A20541357.txt hdfs:///user/csp554-2/
```

```
[hadoop@ip-172-31-35-150 ~]$ hadoop fs -cp s3://a20541357-csp554/A20541357.txt hdfs:///user/csp554-2/
2023-09-15 02:45:20,699 INFO s3n.S3NativeFileSystem: Opening 's3://a20541357-csp554/A20541357.txt' for reading
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -ls /user/csp554-2/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

Found 2 items
-rw-r--r-- 1 hadoop hdfsadmingroup 21 2023-09-15 02:45 /user/csp554-2/A20541357.txt
-rw-r--r-- 1 hadoop hdfsadmingroup 21 2023-09-15 02:36 /user/csp554-2/ankanmazumdar.txt
```

17. (2 points) Execute a command to show the contents of the myid.txt file in the hdfs directory /user/csp554-2. Clue: look up about how to use the "cat" command in the file system shell document. Write down the command you executed and also take a screen snapshot of the listed content of the file and include it in your assignment submission.

```
hadoop dfs -cat /user/csp554-2/A20541357.txt
```

```
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -cat /user/csp554-2/A20541357.txt
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

this is the name file[hadoop@ip-172-31-35-150 ~]$
[hadoop@ip-172-31-35-150 ~]$
```

18. (2 points) Execute a command to remove the myid.txt file in the hdfs directory /user/csp554-2. Clue: look up about how to use the "rm" command in the file system shell document. Write down the

command you executed, then list the content of the /user/csp554-2 HDFS directory and take a screen snapshot of the listed content of the directory and include it in your assignment submission.

```
hadoop dfs -rm -f /user/csp554-2/A20541357.txt
```

```
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -ls /user/csp554-2/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

Found 2 items
-rw-r--r--  1 hadoop hdfsadmingroup    21 2023-09-15 02:45 /user/csp554-2/A20541357.txt
-rw-r--r--  1 hadoop hdfsadmingroup    21 2023-09-15 02:36 /user/csp554-2/ankanmazumdar.txt
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -cat /user/csp554-2/A20541357.txt
```

```
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -rm -f /user/csp554-2/A20541357.txt
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

Deleted /user/csp554-2/A20541357.txt
[hadoop@ip-172-31-35-150 ~]$ hadoop dfs -ls /user/csp554-2/
WARNING: Use of this script to execute dfs is deprecated.
WARNING: Attempting to execute replacement "hdfs dfs" instead.

Found 1 items
-rw-r--r--  1 hadoop hdfsadmingroup    21 2023-09-15 02:36 /user/csp554-2/ankanmazumdar.txt
[hadoop@ip-172-31-35-150 ~]$
```